
CMPE-655 Multiple Processor Systems
Assignment 2
Parallel Implementation of a Ray Tracer

Brandon Key
Submitted: 15 April 2020

Instructor: Dr. Shaaban
TAs: Robert Mason
Adam Smith
Designed by:

1 Abstract

2 Design Methodology and Theory

The given sequential code was the basis for all partitioning methods.

2.1 Static Strips

The first partitioning scheme that was implemented was static vertical strips. Each processor was responsible for rendering an even share of the image as seen in Equation 1.

$$ColumnstoProcess = \frac{ImageWidth}{NumberOfProcessors} \quad (1)$$

Equation 1: Strip Size

Since the calculation for strip size is integer based, there is a remainder. These remaining columns were given to the lowest ranked processors by giving those lowest ranked processors one extra column. Figure 1 shows the division of the image to processors.




Figure 1: Static Vertical Strips

The computation time also needed to be sent to the master along with the pixels. This was done by allocating an extra float to the end of the pixel array and appending the computation

time there. The master would iterate through the slave ranks and accept the packet of pixels. The master would store the computation time if it was the largest one encountered and copy the pixels in the packet into the master packet array.

It should be noted that in order to facilitate vertical strips, the slaves represented the pixels in row major form so that computations happened in a contiguous memory. This at least allowed the pixel array to be in cache more often since there is more locality.

2.2 Static Cycles

The next partition scheme was cyclical assignment. Instead of processors handling a contiguous strip of pixels, a processor handles a few rows of pixels at a time. Figure 2 shows how 4 processors would be responsible for calculations.



Figure 2: Static Horizontal Cycles

It was determined that an iterative approach would best handle the required calculations; each processor would iterate over all rows, and check if it should process that row. Equation 2 shows the check that each processor performs to check if it should calculate the pixel value.

$$rank == \frac{row}{CycleSize} \% Processors \quad (2)$$

Equation 2: Rows Each Processor is Responsible For

An iterative approach simplifies design, and the check for row calculation is much less intensive than the calculation being performed, little performance is lost. This check naturally leads to the lower ranked tasks handling the "leftover" rows. Since the image cannot always be divided evenly by the cycle size, calculations are stopped when the row reaches image edge (via bounds of a for loop).

When the master receives slave results, it must copy the rows from the packet into the pixel array. This was trivial as both the pixel array and the packet array had the same row width. While going through the packet, the master kept track of which row (from the slave perspective) it was on in order to reverse the logic that slave performed. At the end of the packet was the computation time again.

2.3 Static Blocks

The third partitioning scheme is static blocks. The image was broken into (nearly) equal size blocks. The number of blocks was determined by the number of processors available. The width of the image in blocks is the square root of the number of processors. Equation 3a shows the calculation of the width of the image in blocks. Note that the number of processors is limited to perfect squares.

$$\text{Image Size (blocks)} == \text{Image Width (blocks)} * \text{Image Height (blocks)} \quad (3a)$$

$$\text{Image Width (blocks)} = \sqrt{\text{Image Size}} \quad (3b)$$

$$\text{Image Height (blocks)} = \sqrt{\text{Image Size}} \quad (3c)$$

Equation 3a: Rows Each Processor is Responsible For

Figure 3 shows the division of the image between processors.



Figure 3: Static Blocks

The figure shows that blocks on the end are not the same size as the blocks at the beginning. In order to accommodate the remainder of the pixels that results from dividing the image size by

$\sqrt{Processors}$, the remaining pixels are given to the last n blocks. By putting the larger blocks closer to the end, it gives the master time to communicate and handle the data from the lower ranked slaves, which hopefully finisher quicker.

To facilitate computation, a struct the represents a block was created. It specifies the coordinates of block in terms of blocks and pixels. This allows easily determining the appropriate packet size as well as which pixels the tasks should process. This helped both slave and master as well as cleaned up the code.

Like with other static partitioning schemes, the master would receive slave packets in order of their rank. Again, the packet would

2.4 Dynamic Blocks




Figure 4: Dynamic Blocks

3 Results and Analysis

3.1 Sequential

The sequential code took 112.94 seconds to parse the simple image and 3585.36 seconds to parse the complex image.

3.2 Static Strips

Table 1: Performance of Vertical Strips with Simple Image

Static Strip Allocation

Processes	Strips	Exec. Time (s)	Speedup	C-to-C ratio
1 (srun -n 1)	1	175.619535	0.643	0
2 (srun -n 2)	2	121.061101	0.933	0.000505
4 (srun -n 4)	4	82.225865	1.374	0.002754
9 (srun -n 9)	9	41.130541	2.746	0.003184
16 (srun -n 16)	16	26.236589	4.305	0.028886
20 (srun -n 20)	20	22.363491	5.050	0.118175
25 (srun -n 25)	25	18.112424	6.235	0.097723
36 (srun -n 36)	36	12.570574	8.984	0.093093
49 (srun -n 49)	49	10.159625	11.117	0.19525
55 (srun -n 55)	55	9.505927	11.881	0.188241
64 (srun -n 64)	64	8.895179	12.697	0.368489

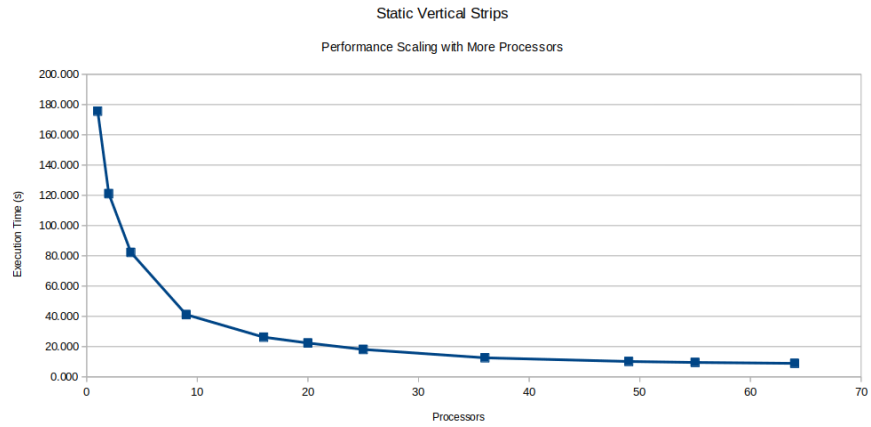


Figure 5: Performance Scaling of Static Vertical Strips

3.3 Static Cycles

Table 2: Performance of Horizontal Cycles with Simple Image as Cycle Size Varies

Static Cyclical Allocation				
Processes	Strip Height (pixels)	Exec. Time (s)	Speedup	C-to-C ratio
4 (srun -n 4)	1	47.654	2.370	0.011687
4 (srun -n 4)	5	47.649	2.370	0.004629
4 (srun -n 4)	10	47.672	2.369	0.011275
4 (srun -n 4)	20	47.873	2.359	0.00514
4 (srun -n 4)	80	48.269	2.340	0.01162
4 (srun -n 4)	320	49.859	2.265	0.005853
4 (srun -n 4)	640	52.365	2.157	0.005939
4 (srun -n 4)	1280	82.994	1.361	0.002594
9 (srun -n 9)	1	22.686	4.978	0.078114
9 (srun -n 9)	5	21.511	5.250	0.005657
9 (srun -n 9)	10	22.969	4.917	0.083974
9 (srun -n 9)	20	21.982	5.138	0.017573
9 (srun -n 9)	40	23.232	4.861	0.070674
9 (srun -n 9)	160	25.078	4.504	0.016896
9 (srun -n 9)	350	28.399	3.977	0.016553
9 (srun -n 9)	650	48.166	2.345	0.005191
16 (srun -n 16)	1	14.010	8.061	0.179885
16 (srun -n 16)	5	13.371	8.446	0.124782
16 (srun -n 16)	10	14.056	8.035	0.177079
16 (srun -n 16)	20	13.519	8.354	0.115796
16 (srun -n 16)	50	14.400	7.843	0.132413
16 (srun -n 16)	100	16.255	6.948	0.135441
16 (srun -n 16)	250	21.617	5.225	0.077393
16 (srun -n 16)	400	34.928	3.234	0.089348

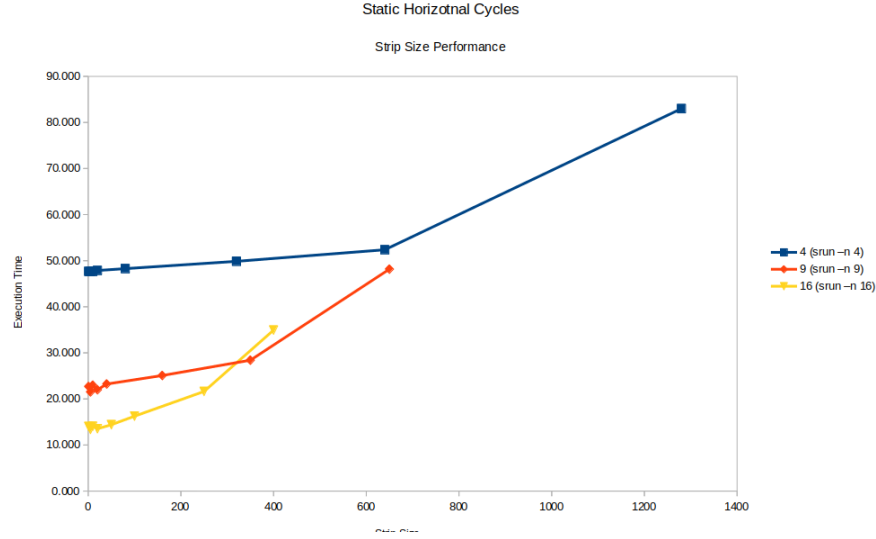


Figure 6: Performance Scaling of Cycle Size

Table 3: Performance of Horizontal Cycles with Simple Image as Number of Processors Varies

Static Cyclical Allocation				
Processes	Strip Height (pixels)	Exec. Time (s)	Speedup	C-to-C ratio
1 (srun -n 1)	27	173.187	0.652	0
2 (srun -n 2)	27	89.103	1.268	0.002158
4 (srun -n 4)	27	48.128	2.347	0.005378
9 (srun -n 9)	27	24.305	4.647	0.131897
16 (srun -n 16)	27	13.658	8.269	0.120874
20 (srun -n 20)	27	14.311	7.892	0.394911
25 (srun -n 25)	27	10.655	10.599	0.328285
36 (srun -n 36)	27	10.214	11.057	0.723326
49 (srun -n 49)	27	8.846	12.768	1.087796
55 (srun -n 55)	27	12.878	8.770	2.074718
64 (srun -n 64)	27	7.341	15.384	0.927029




Figure 7: Performance Scaling with Increasing Number of Processors

3.4 Static Blocks

Table 4: Performance of Blocks with Simple Image

Static Block Allocation				
Processes	Blocks	Exec. Time (s)	Speedup	C-to-C ratio
1 (srun -n 1)	1	179.020	0.631	0
4 (srun -n 4)	4	74.761	1.511	0.01028
9 (srun -n 9)	9	55.282	2.043	0.073443
16 (srun -n 16)	16	35.953	3.141	0.027227
25 (srun -n 25)	25	22.797	4.954	0.139857
36 (srun -n 36)	36	21.581	5.233	0.166538
49 (srun -n 49)	49	18.009	6.271	0.334002
64 (srun -n 64)	64	15.210	7.425	0.267504

3.5 Dynamic Blocks

Table 5: Performance of Blocks - Simple Image with Square Blocks

Dynamic Allocation				
Processes	Block Size (rows x cols)	Exec. Time (s)	Speedup	C-to-C Ratio
9 (srun -n 9)	1x1	38.433	2.939	7.30E-07
9 (srun -n 9)	15x15	25.908	4.359	7.94E-07
9 (srun -n 9)	25x25	26.083	4.330	9.49E-07
9 (srun -n 9)	50x50	23.540	4.798	1.10E-06
9 (srun -n 9)	75x75	24.575	4.596	1.10E-06
9 (srun -n 9)	100x100	23.523	4.801	1.00E-06
16 (srun -n 16)	1x1	25.339	4.457	1.20E-07
16 (srun -n 16)	15x15	14.009	8.062	1.37E-06
16 (srun -n 16)	25x25	14.461	7.810	1.98E-06
16 (srun -n 16)	50x50	14.185	7.962	1.45E-06
16 (srun -n 16)	75x75	13.549	8.336	2.02E-06
16 (srun -n 16)	100x100	12.938	8.729	1.54E-06

Table 6: Performance of Blocks - Simple Image with Oblong Blocks

Dynamic Allocation				
Processes	Block Size (rows x cols)	Exec. Time (s)	Speedup	C-to-C Ratio
16 (srun -n 16)	1x11	31.751	3.557	1.04E-06
16 (srun -n 16)	10x1	37.412	3.019	5.17E-07
16 (srun -n 16)	13x7	15.367	7.350	2.26E-06
16 (srun -n 16)	29x5	15.644	7.219	1.19E-06
16 (srun -n 16)	110x9	13.726	8.228	1.99E-06
16 (srun -n 16)	1x70	16.945	6.665	1.13E-06
36 (srun -n 36)	1x11	46.007	2.455	3.24E-07
36 (srun -n 36)	10x1	49.738	2.271	3.08E-07
36 (srun -n 36)	13x7	8.391	13.460	3.18E-06
36 (srun -n 36)	29x5	8.282	13.636	2.33E-06
36 (srun -n 36)	110x9	6.521	17.320	4.19E-06
36 (srun -n 36)	1x70	10.135	11.144	2.23E-06

Table 7: Performance of Blocks - Complex Image with Square Blocks

Dynamic Allocation				
Processes	rows x cols	Exec. Time (s)	Speedup	C-to-C Ratio
9 (srun -n 9)	1x1	752.421	4.765	3.51E-08
9 (srun -n 9)	15x15	742.651	4.828	2.38E-08
9 (srun -n 9)	25x25	764.933	4.687	2.05E-08
9(srun -n 9)	50x50	810.189	4.425	1.91E-08
9 (srun -n 9)	75x75	796.143	4.503	2.05E-08
9 (srun -n 9)	100x100	856.743	4.185	1.96E-08
16 (srun -n 16)	1x1	827.368	4.333	2.48E-08
16 (srun -n 16)	15x15	405.382	8.844	3.36E-08
16 (srun -n 16)	25x25	437.604	8.193	3.59E-08
16 (srun -n 16)	50x50	465.604	7.700	3.30E-08
16 (srun -n 16)	75x75	467.381	7.671	3.95E-08
16 (srun -n 16)	100x100	483.416	7.417	3.25E-08

Table 8: Performance of Blocks - Complex Image with Oblong Blocks

Dynamic Allocation				
Processes	rows x cols	Exec. Time (s)	Speedup	C-to-C Ratio
16 (srun -n 16)	1x11	413.590	8.669	4.87E-08
16 (srun -n 16)	10x1	411.836	8.706	6.49E-08
16 (srun -n 16)	13x7	398.067	9.007	3.97E-08
16 (srun -n 16)	29x5	400.791	8.946	3.91E-08
16 (srun -n 16)	110x9	433.408	8.272	4.08E-08
16 (srun -n 16)	1x70	398.639	8.994	2.24E-07
36 (srun -n 36)	1x11	200.745	17.860	9.14E-08
36 (srun -n 36)	10x1	240.474	14.910	8.40E-08
36 (srun -n 36)	13x7	176.629	20.299	9.01E-08
36 (srun -n 36)	29x5	181.390	19.766	9.16E-08
36 (srun -n 36)	110x9	214.380	16.724	9.30E-08
36 (srun -n 36)	1x70	172.041	20.840	1.06E-07

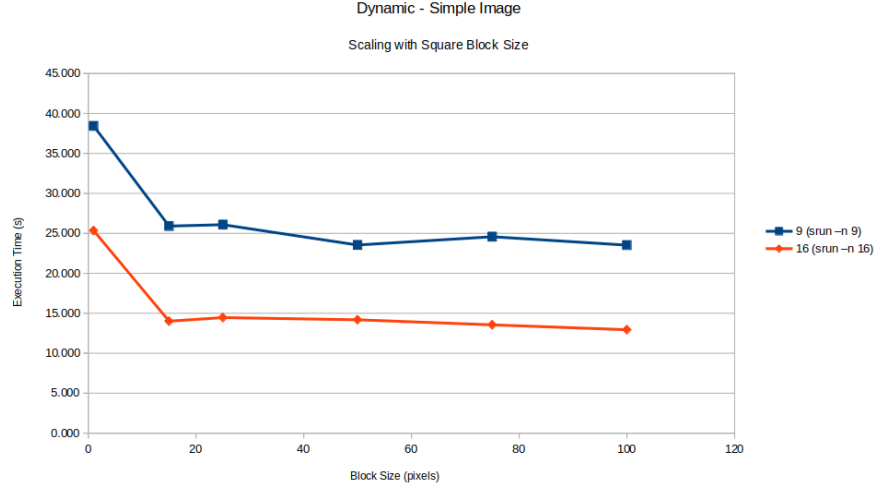


Figure 8: Performance Scaling of Dynamic Mapping Processing Simple Image




Figure 9: Performance Scaling of Dynamic Mapping Processing Complex Image

4 Conclusion

Since dynamic allocation was the last implementation, many lessons were learned during its implementation that could be applied to the static partitioning methods. In all static methods, the master waited for slave results in order of their rank. This causes blocking behavior if a slave with lower rank takes longer to finish processing than a slave with higher rank. To fix this, the master should accept results from any slave, and the slave number should have been included in the packet. Additionally, the master should be given the least amount of processing so it has time to process the slave results.

Dynamic cycles probably the best

The static partitioning schemes try to balance load by giving each processor equal amounts of work. This works well if processors are homogeneous and only the ray-tracing program is running on the processors. If one processor is significantly faster (or less loaded) than the rest, it will sit idle, the converse is true too: a slower processor (or a more highly loaded one) will hold up the rest

of the processors. Dynamic scheduling may have extra overhead in terms of the master not doing work and additional communication, but the more even load balancing tends to hide this overhead. The greater the difference in processors, the greater the advantage that dynamic scheduling will have.

5 Acknowledgments